


Introduction to Matlab

**Prepared by Ohad Shamir, based
on slides by Yonatan Amit**

- 
- This presentation is intended to provide a quick and dirty introduction to Matlab, focusing on stuff you'll need for the programming mini-project.
 - There are tons of other things you can do with Matlab.
 - If you have a question which is not answered here, see Matlab's help, google it, or email me.



Why Matlab?

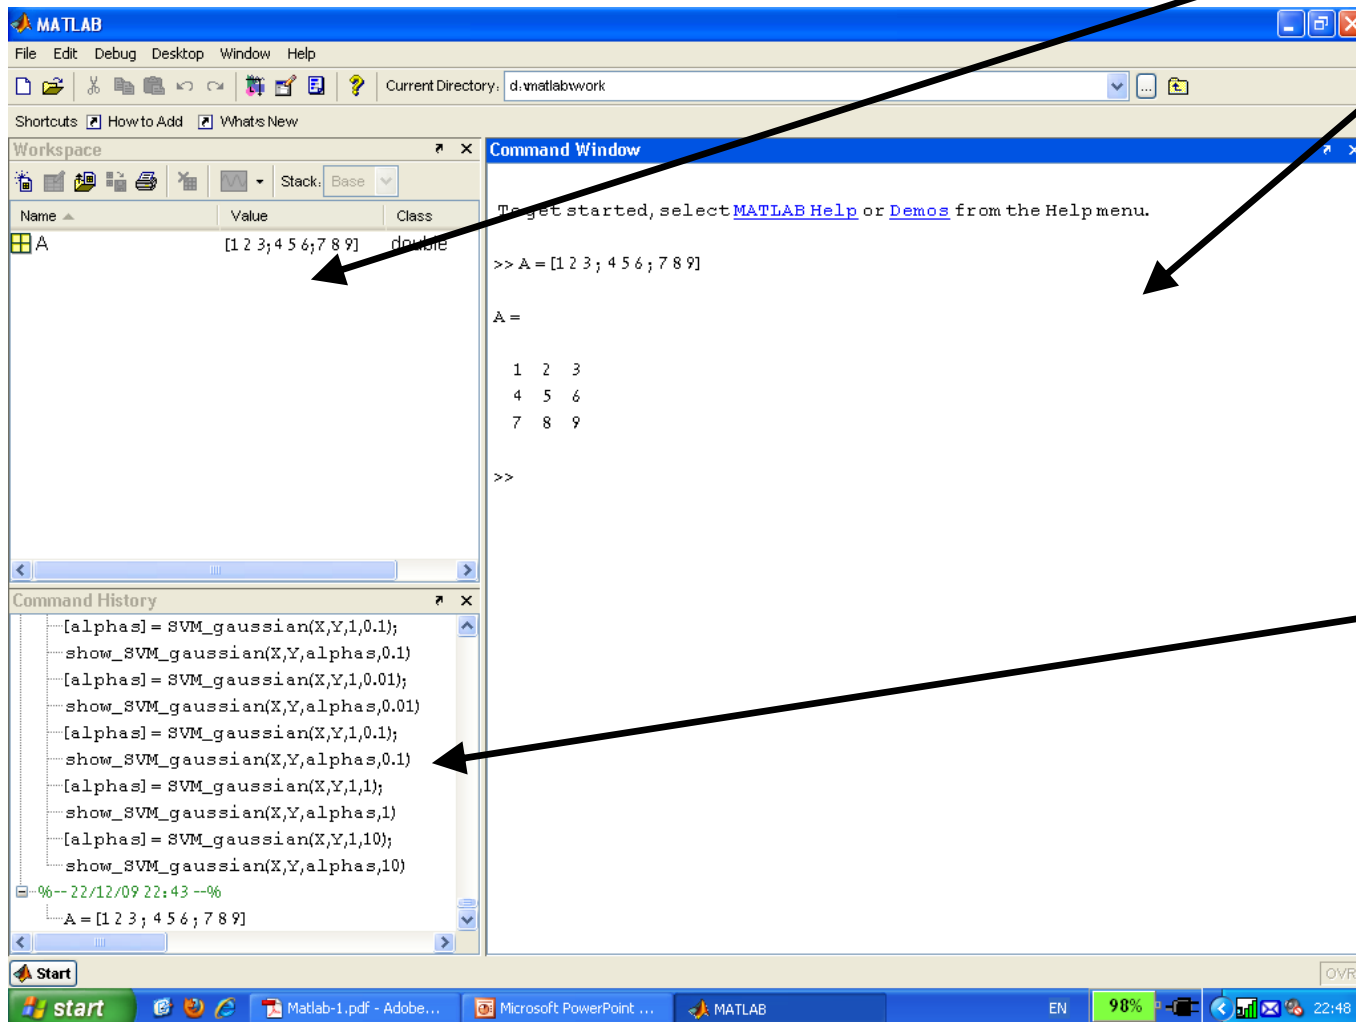
- Matlab is a programming tool for the lazy.
- Used mainly for numerical computation.
 - Basic data structure is a matrix.
 - Has tons of functions for mathematical calculations.
- Generally easy and forgiving (no memory management; programs are interpreted rather than compiled...)
- Excellent for trying out algorithms based on numerical computations (such as those taught in the course)
 - For full-scale, speed-optimized implementations, other programming languages are often used (such as C++).
- Matlab is installed on all public computers on campus.

Welcome to Matlab

Variables

Work area
(command
line
interface)

History



Basic Workflow

- Commands are executed from the command line interface.
- Basically, an elaborate calculator.
 - Try typing `1+1 <enter>`.
 - Try `log(sqrt(3)+4^2) <enter>`
 - Can also store variables! Try `x=3 <enter>`,
`sin(x) <enter>`

Basic Workflow

- (Almost) anything in Matlab is a matrix.
- Scalars (like x in the previous slide) are just a 1×1 matrix.
- Try the following and see what happens:
 - $A = [1 \ 2 \ 3 \ ; \ 4 \ 5 \ 6 \ ; \ 7 \ 8 \ 9]$
 - $B = [4; 5; 6]$
 - $w = A * B$ (multiplies the matrices)
 - $w * A$ (should return an error..)
 - w' (this transposes w)
 - $w' * A$ (now it works...)

Easy Ways to Create Matrices and Vectors

■ Try the following:

- `A=1 : 3`
- `A= (1 : 3) '`
- `A = 1 : 2 : 6` (go from 1 till 6 in steps of 2)
- `A = 5 : -2 : 0` (go from 5 down to 0 in steps of -2)
- `A = zeros (5, 2)` (create a zero matrix with 5 rows and 2 columns)
- `A = ones (3, 4)` (create a matrix of ones with 3 rows and 4 columns)
- `A = eye (5)` (create an identity matrix of size 5*5).

Hiding the Output

- Often, we want to do a calculation without seeing the output explicitly.
 - Try `A=ones(500,500)`. This should look ugly..
 - Tip: press `CTRL+<c>` to halt a computation which takes too long.
 - Now try `A=ones(500,500);`
 - The matrix was created and stored as A, without us seeing the matrix explicitly.

Accessing a matrix

■ Given a matrix A:

- We can access its i, j element: $A(i, j)$
- We can access its i 'th column: $A(:, i)$
- We can access its last column: $A(:, \text{end})$
- We can access the last 10 rows:

$A(\text{end}-10:\text{end}, :)$

- It is also useful for assignments:

$A(1, [3 \ 4]) = [5 \ 6]$

- And so forth...

What else we can do with matrices

- Erase them: `clear A`
- Erase all of them: `clear`
- Save them: `save filename A`
 - Data is saved into a `.mat` file
- Load them: `load filename`
- List them: `who` or `whos`

Operations on Matrices

■ Scalar operations

- $A * 5$ (multiplies all entries by 5)
- $A + 5$ (add 5 to all entries)

■ Matrix operations

- $A * B, A + B$
- A'

■ Element-wise operations

- $A .* B$
- $A .^ B$

Matrix dimensions
must agree





Scripts and Functions

- What if we want to do a complicated calculation? (e.g. run a perceptron algorithm on some data...)
- Matlab provides two mechanisms to encapsulate code: scripts and functions.
- In both cases, you write code as a separate file (called an m-file), and then invoke the script/function from the command line.

Scripts

- Just a sequence of commands.
- Go to File->New->M-file. This opens the m-file editor.
- Type (for example):

```
A = [1 2 3];
```

```
B = [4 5 6];
```

```
A.*B
```

- Save this file as `xxx.m` in Matlab's current directory.
- Now, go to the command prompt, and type `xxx`.

Functions

- Also written as an m-file.
- Must have input and an output.
- M-file must begin with a declaration such as:
`[a,b] = my_function(c,d,e)`
- File name should be the same as function name.
- The file may contain other private functions (not needed for the miniproject...)
- The function can be invoked from the command line, just like any other function.



If...

- Create some scalar `a`, and try to run the following script:

```
if (a==5)
    display('a is just right');
else
    display('a is not right');
end
```



If...

- Create some scalar a , and try to run the following script:

```
if (a>5)
    display('a is too big');
elseif (a<5)
    display('a is too small');
else
    display('a is just right');
end
```

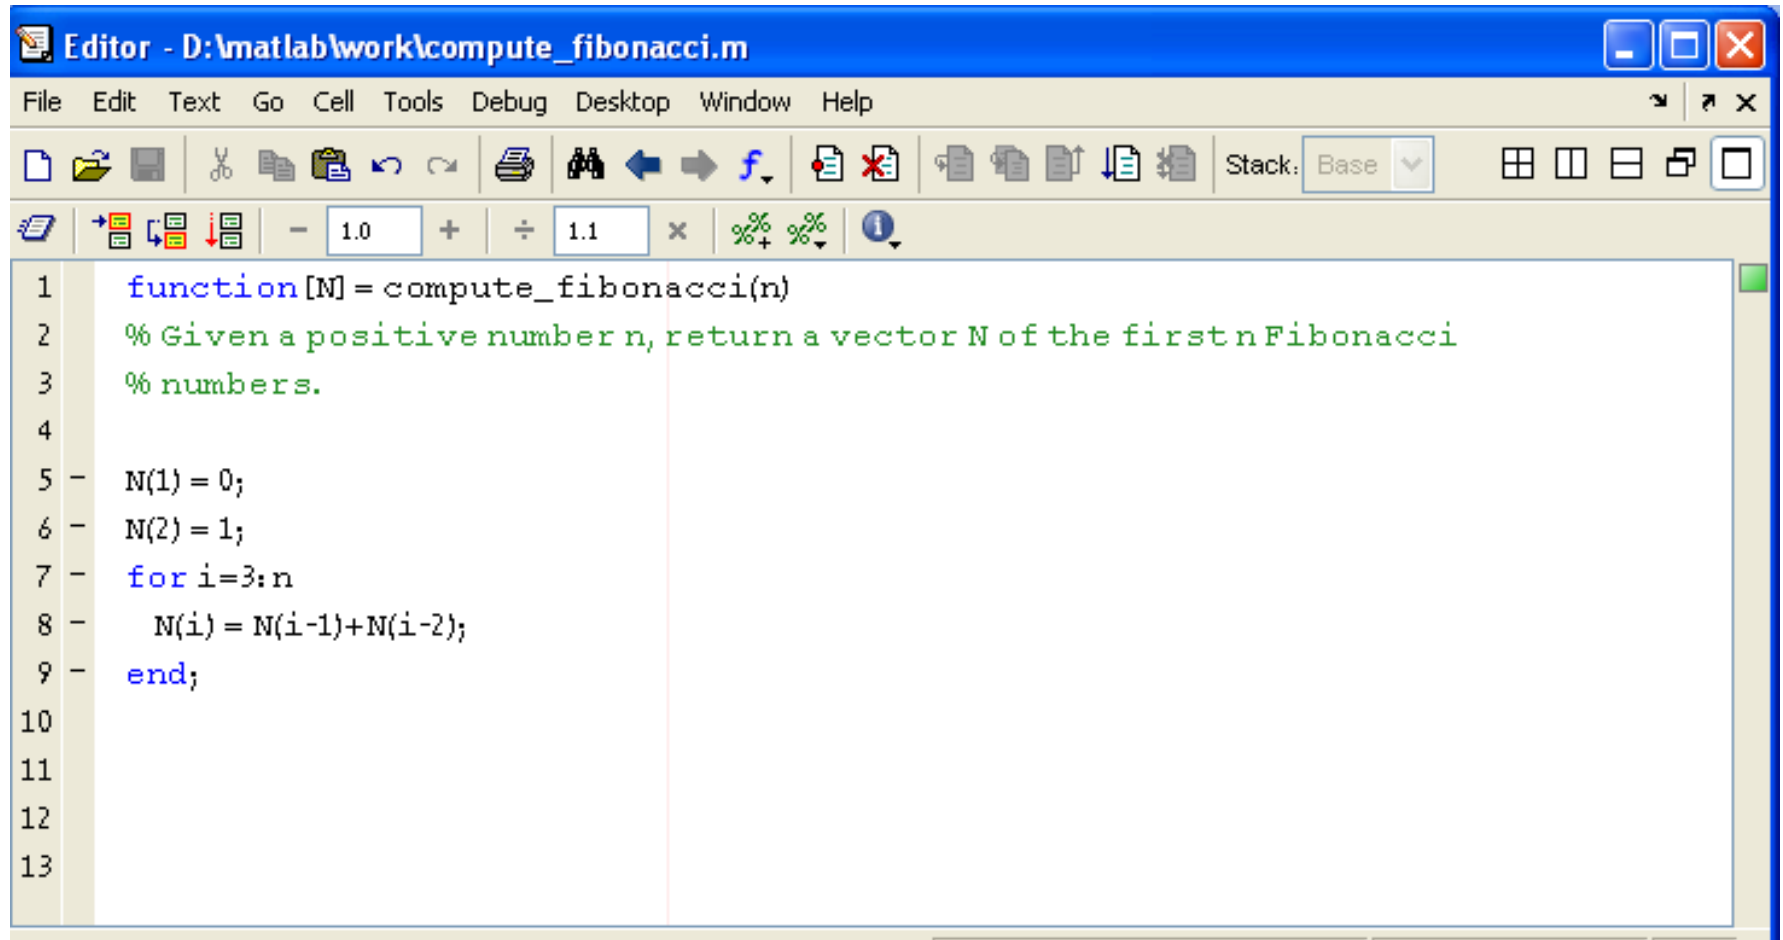



Looping

```
For i=[1,6:8]  
    i+3  
end
```

```
i=0;  
while i<3  
    i = i+1;  
    disp('Hello');  
end
```

Example



The image shows a MATLAB Editor window titled "Editor - D:\matlab\work\compute_fibonacci.m". The window has a menu bar (File, Edit, Text, Go, Cell, Tools, Debug, Desktop, Window, Help) and a toolbar with various icons for file operations, editing, and execution. Below the toolbar is a stack of variables, currently showing "Base". The main editing area contains the following MATLAB code:

```
1 function [N] = compute_fibonacci(n)
2 % Given a positive number n, return a vector N of the first n Fibonacci
3 % numbers.
4
5 - N(1) = 0;
6 - N(2) = 1;
7 - for i=3:n
8 -     N(i) = N(i-1)+N(i-2);
9 - end;
```

The code defines a function `compute_fibonacci` that takes an input `n` and returns a vector `N` containing the first `n` Fibonacci numbers. The function initializes `N(1) = 0` and `N(2) = 1`, then uses a `for` loop to calculate the subsequent Fibonacci numbers up to `n`.

Example

The image shows the MATLAB software interface. The top menu bar includes File, Edit, Debug, Desktop, Window, and Help. Below the menu bar is a toolbar with various icons. The current directory is set to d:\matlab\work. The workspace window displays three variables: A, a, and ans, all of type double. The command window shows the execution of the compute_fibonacci function. The command history window shows a list of commands executed, including SVM_gaussian and compute_fibonacci.

Workspace

| Name | Value | Class |
|------|-----------------------|--------|
| A | [1 2 3; 4 5 6; 7 8 9] | double |
| a | [0 1 1 2 3 5] | double |
| ans | [0 1 1 4 9 25] | double |

Command Window

```
>> compute_fibonacci(6)

ans =

    0    1    1    2    3    5

>> a = compute_fibonacci(6);
>> a.^2

ans =

    0    1    1    4    9   25

>>
```

Command History

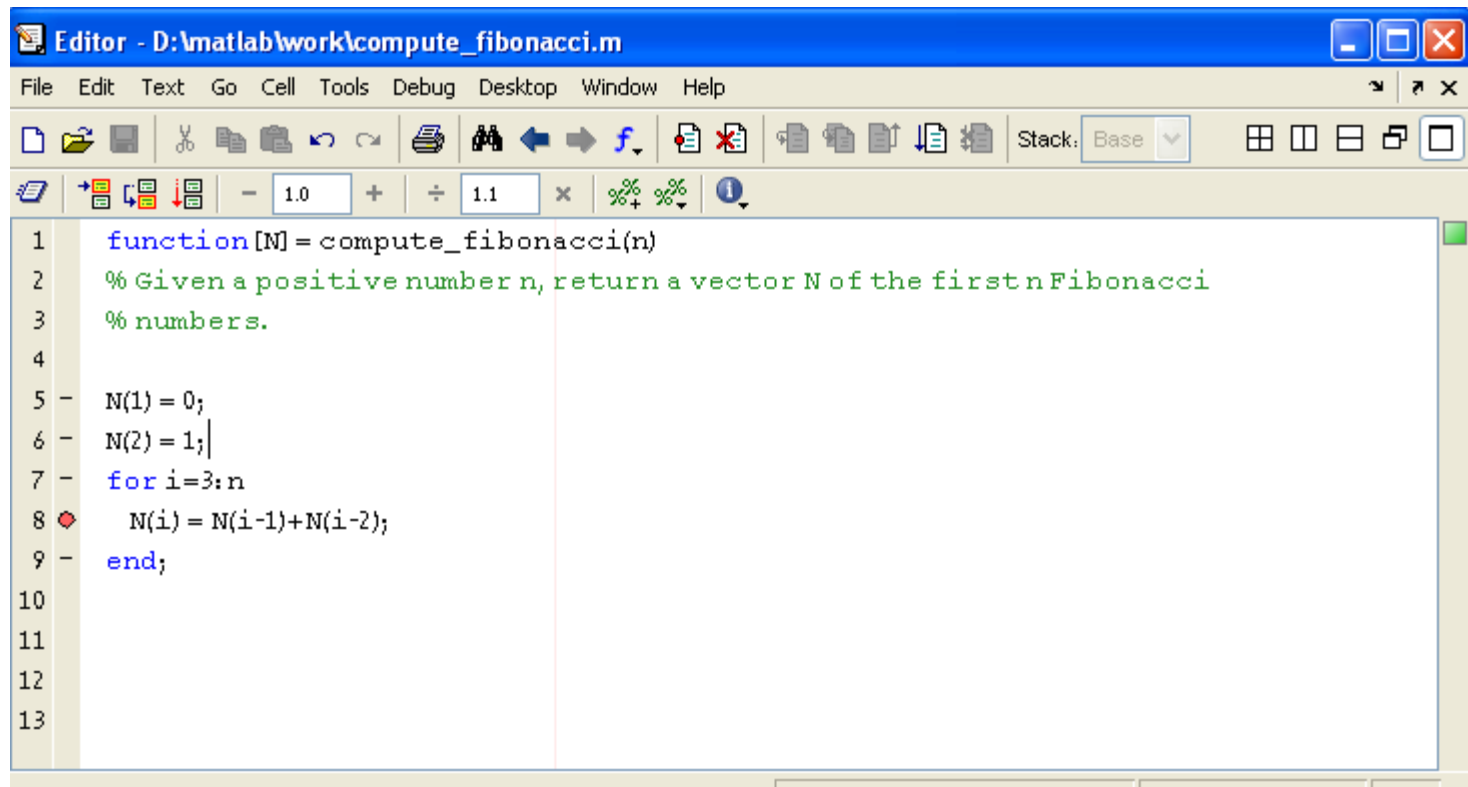
```
[alphas] = SVM_gaussian(X,Y,1,1);
show_SVM_gaussian(X,Y,alphas,1)
[alphas] = SVM_gaussian(X,Y,1,10);
show_SVM_gaussian(X,Y,alphas,10)
--22/12/09 22:43 --%
A = [1 2 3; 4 5 6; 7 8 9]
5:-2:0
compute_fibonacci(4)
compute_fibonacci(4)
compute_fibonacci(6)
a = compute_fibonacci(6);
a.^2
```

Taskbar

Start button, taskbar with open applications: Matlab-1.pdf..., Microsoft Po..., MATLAB, Editor - D:\m..., Fibonacci num..., EN, 98% battery, 23:37.

The Matlab Debugger

- Edit a file and mark a line by clicking on the gray area to the left of the line.





The Matlab Debugger

- When you run the function, the execution will stop when the breakpoint is reached, and control is returned to the command line.
- You can then check the variables, add/remove breakpoints etc.
- Use the debug menu for continuing (continue as normal, step one command ahead etc...)
- Can also add conditions to the breakpoints.



Programming Matlab the right way

- Matlab is very fast in doing calculations over matrices.
- It is *very* slow in doing loops.
- Always try to avoid loops if possible, by doing vector/matrix calculations.

Programming Matlab the right way

- Suppose you want to calculate the norm of a kernel classifier, given weights $\alpha_1, \dots, \alpha_m$ and the training instances x_1, \dots, x_m . Recall that:

$$\begin{aligned}\|w\|^2 &= \left\langle \sum_{i=1}^m \alpha_i \Psi(x_i), \sum_{i=1}^m \alpha_i \Psi(x_i) \right\rangle \\ &= \sum_{i,j=1}^m \alpha_i \alpha_j \langle \Psi(x_i), \Psi(x_j) \rangle \\ &= \sum_{i,j=1}^m \alpha_i \alpha_j K(x_i, x_j)\end{aligned}$$

Programming Matlab the right way

- So, if we have a matrix G where row i and column j contain $K(x_i, x_j)$, and a vector `alphas` where `alphas(i)` is α_i , the naïve way to calculate the norm is:

```
norm2 = 0;
for i=1:m
    for j=1:m
        norm2 = norm2+alphas(i)*alphas(j)*G(i,j);
    end
end
classifier_norm = sqrt(norm2);
```


Programming Matlab the right way

- The smart way is to do the following (try to understand why it is equivalent)

```
classifier_norm = sqrt(alphas'*G*alphas);
```

- On my laptop, on a 1000×1000 matrix G , this implementation runs 435 times faster than the naïve implementation!

A list of useful functions

- `repmat` : replicate a vector/matrix, e.g.

```
>> repmat([1 2 3], 3, 1)
```

```
ans =
```

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 1 | 2 | 3 |
| 1 | 2 | 3 |



A list of useful functions

- `size(X, 1)` : number of rows of matrix X .
- `size(X, 2)` : number of columns of matrix X .
- `sum(a)` : sum the elements of a vector a .
- `sum(X, 1)` : sum the rows of matrix X .
- `sum(X, 2)` : sum the columns of matrix X .

A list of useful functions

- `rand`: Return a random number uniformly distributed on $[0,1]$.
- `randi(n)`: Return a random number uniformly distributed on $\{0,1,\dots,n\}$.
- `randperm(m)`: return a random permutation of the numbers $1,2,\dots,m$
- `floor(x)`: return the largest integer smaller than a number x .
- `[val, ind] = min(a)`: return the smallest element val in vector a , so that $a(ind)$ is equal to val .
- `find(a>5)`: return the indices of the entries in vector a which are larger than 5 (for example...)